

Spring Boot Transaction Management - Comprehensive Notes

1. Introduction to Transactions in Spring Boot

What are Transactions?

- **Definition:** A transaction is a sequence of database operations that are treated as a single unit of work
- **Purpose:** Ensures data integrity and consistency in database operations
- **Spring Boot Integration:** Provides seamless integration with Hibernate/JPA through declarative transaction management

Why Use Spring Transactions?

- **Simplified Management:** No need for manual transaction handling with Hibernate
 - **Declarative Approach:** Use annotations instead of programmatic transaction code
 - **Automatic Rollback:** Handles exceptions and rollbacks automatically
 - **Connection Management:** Manages database connections efficiently
-

2. @Transactional Annotation

Basic Usage

- **Primary Annotation:** `@Transactional` is the main annotation for transaction management
- **Declarative Style:** Eliminates boilerplate transaction management code
- **Automatic Proxy Creation:** Spring creates proxies to handle transaction lifecycle

How @Transactional Works Internally

```
java

// Spring wraps your method with transaction management code
createTransactionIfNecessary();
try {
    yourMethod(); // Your actual business logic
    commitTransactionAfterReturning();
} catch (Exception exception) {
    rollbackTransactionAfterThrowing();
    throw exception;
}
```

Application Levels (Priority Order - Lowest to Highest)

1. **Interface Level**: Applied to entire interface
2. **Superclass Level**: Applied to parent class
3. **Class Level**: Applied to entire class
4. **Interface Method Level**: Applied to specific interface method
5. **Superclass Method Level**: Applied to parent class method
6. **Method Level**: Applied to specific method (Highest Priority)

Example Implementation

```
java

@Transactional(readOnly = true) // Class Level - default read-only
public class EmployeeService {

    @Autowired
    private EmployeeRepository employeeRepo;

    @Transactional // Method Level - overrides class level settings
    public Employee addEmployee(Employee employee) {
        Employee savedEmployee = employeeRepo.save(employee);
        // Additional operations within same transaction
        return savedEmployee;
    }
}
```

3. ACID Properties of Transactions

A - Atomicity

- **All or Nothing Principle**: Either all operations succeed or all fail
- **Complete Execution**: All operations within transaction are executed successfully
- **Rollback Capability**: If any operation fails, entire transaction is rolled back
- **Data Integrity**: Prevents partial updates that could corrupt data

C - Consistency

- **State Transition**: Moves database from one consistent state to another
- **Constraint Validation**: All database constraints must be satisfied
- **Business Rules**: Ensures business logic rules are maintained
- **Data Validity**: Modified data passes all validation checks

I - Isolation

- **Transaction Separation:** Changes are invisible to other transactions until commit
- **Concurrent Access:** Multiple transactions can run simultaneously without interference
- **Data Protection:** Prevents dirty reads and inconsistent data access
- **Isolation Levels:** Different levels control visibility of uncommitted changes

D - Durability

- **Persistent Storage:** Committed changes are permanently saved
 - **System Failure Recovery:** Data survives system crashes and restarts
 - **Storage Guarantee:** Changes are written to non-volatile storage
 - **Long-term Reliability:** Data remains intact over time
-

4. Transaction Propagation Types

REQUIRED (Default)

- **Behavior:** Join existing transaction or create new one
- **Usage:** Most common propagation type
- **Characteristics:**
 - Checks for active transaction
 - Creates new transaction if none exists
 - Joins existing transaction if available
 - Exception rolls back entire transaction (parent + child)

java

```
@Transactional(propagation = Propagation.REQUIRED)
public void methodWithRequired() {
    // Business Logic here
}
```

REQUIRES_NEW

- **Behavior:** Always create a completely new transaction
- **Usage:** When you need independent transaction handling
- **Characteristics:**
 - Suspends current transaction if exists
 - Creates new physical transaction with new database connection

- Outer transaction resumes after inner completes
- Independent rollback behavior

java

```
@Transactional(propagation = Propagation.REQUIRES_NEW)
public void methodWithRequiresNew() {
    // Runs in completely separate transaction
}
```

NESTED

- **Behavior:** Uses savepoints for nested operations
- **Usage:** When you need partial rollback capability
- **Characteristics:**
 - Creates savepoint if transaction exists
 - Rollback only to savepoint on exception
 - Works like REQUIRED if no active transaction
 - Allows partial transaction recovery

java

```
@Transactional(propagation = Propagation.NESTED)
public void methodWithNested() {
    // Can rollback to savepoint if needed
}
```

MANDATORY

- **Behavior:** Requires existing transaction, throws exception if none
- **Usage:** When method must run within existing transaction context
- **Characteristics:**
 - Never creates new transaction
 - Throws exception if no active transaction
 - Joins existing transaction
 - Rollback affects entire transaction chain

```
java

@Transactional(propagation = Propagation.MANDATORY)
public void methodWithMandatory() {
    // Must be called within existing transaction
}
```

NEVER

- **Behavior:** Throws exception if transaction exists
- **Usage:** When method must run outside transaction context
- **Characteristics:**
 - Prohibits transaction execution
 - Throws IllegalTransactionStateException if transaction exists
 - Internal methods can still create transactions
 - Useful for read-only operations that don't need transactions

```
java

@Transactional(propagation = Propagation.NEVER)
public void methodWithNever() {
    // Must run without any transaction
}
```

NOT_SUPPORTED

- **Behavior:** Suspends existing transaction and runs without transaction
- **Usage:** When you need to temporarily escape transaction context
- **Characteristics:**
 - Suspends current transaction if exists
 - Executes without transaction
 - Resumes original transaction after completion
 - Useful for operations that shouldn't be part of transaction

```
java

@Transactional(propagation = Propagation.NOT_SUPPORTED)
public void methodWithNotSupported() {
    // Runs outside transaction context
}
```

SUPPORTS

- **Behavior:** Flexible - works with or without transaction
- **Usage:** When transaction is optional
- **Characteristics:**
 - Joins existing transaction if available
 - Runs without transaction if none exists
 - Most flexible propagation type
 - Adapts to calling context

java

```
@Transactional(propagation = Propagation.SUPPORTS)
public void methodWithSupports() {
    // Works in any context
}
```

5. Transaction Isolation Levels

Overview

- **Purpose:** Control degree of isolation between concurrent transactions
- **Trade-off:** Balance between consistency and performance
- **Configuration:** Set via `@Transactional(isolation = Isolation.LEVEL)`

Read Phenomena

Dirty Read

- **Definition:** Reading uncommitted data from another transaction
- **Problem:** Data might be rolled back, making read invalid
- **Example:** T1 modifies data, T2 reads it, T1 rolls back
- **Risk:** Working with data that never officially existed

Non-Repeatable Read

- **Definition:** Same query returns different results within same transaction
- **Problem:** Data changes between reads in same transaction
- **Example:** T1 reads data, T2 modifies and commits, T1 reads again with different result
- **Risk:** Inconsistent data within single transaction

Phantom Read

- **Definition:** Same query returns different number of rows

- **Problem:** New rows appear/disappear between reads
- **Example:** T1 counts records, T2 inserts new record, T1 counts again with different total
- **Risk:** Set-based operations become unreliable

Isolation Levels

READ_UNCOMMITTED

- **Allows:** All read phenomena (Dirty, Non-repeatable, Phantom)
- **Performance:** Highest
- **Consistency:** Lowest
- **Use Case:** When performance is critical and consistency is less important

READ_COMMITTED (Default in most databases)

- **Prevents:** Dirty reads
- **Allows:** Non-repeatable reads, Phantom reads
- **Performance:** High
- **Consistency:** Moderate
- **Use Case:** Most common business applications

REPEATABLE_READ

- **Prevents:** Dirty reads, Non-repeatable reads
- **Allows:** Phantom reads
- **Performance:** Moderate
- **Consistency:** High
- **Use Case:** When consistent reads are important

SERIALIZABLE

- **Prevents:** All read phenomena
- **Allows:** None
- **Performance:** Lowest
- **Consistency:** Highest
- **Use Case:** Critical financial or audit operations

6. @Transactional Rules and Best Practices

Method Requirements

- **Visibility:** Must be applied only to `public` methods
- **Invocation:** Method must be called from outside the bean (external invocation)
- **Proxy Limitation:** Internal method calls bypass proxy and ignore `@Transactional`
- **Self-Invocation:** Calling `@Transactional` method from same class won't work

Best Practices

1. **Service Layer:** Apply transactions at service layer, not repository layer
2. **Method Granularity:** Keep transactional methods focused and small
3. **Exception Handling:** Be aware of checked vs unchecked exception rollback behavior
4. **Read-Only Optimization:** Use `readOnly = true` for query-only operations
5. **Timeout Settings:** Set appropriate timeout values for long-running operations
6. **Isolation Selection:** Choose appropriate isolation level based on business needs

Common Configuration Options

```
java
@Transactional(
    propagation = Propagation.REQUIRED,
    isolation = Isolation.READ_COMMITTED,
    timeout = 30,
    readOnly = false,
    rollbackFor = {Exception.class},
    noRollbackFor = {BusinessException.class}
)
```

Performance Considerations

- **Connection Pooling:** Transactions hold database connections
 - **Lock Duration:** Longer transactions increase lock contention
 - **Resource Usage:** Consider memory and CPU impact
 - **Batch Operations:** Use appropriate batch sizes for bulk operations
-

7. Advanced Transaction Concepts

Rollback Behavior

- **Default:** Rollback on unchecked exceptions (`RuntimeException` and `Error`)
- **Checked Exceptions:** Do not trigger rollback by default
- **Custom Configuration:** Use `rollbackFor` and `noRollbackFor` attributes

- **Manual Rollback:** Use `TransactionStatus.setRollbackOnly()`

Transaction Synchronization

- **Before Commit:** Execute code before transaction commits
- **After Commit:** Execute code after successful commit
- **After Rollback:** Execute code after rollback
- **Use Cases:** Cache invalidation, event publishing, cleanup operations

Testing Transactions

- **@Transactional in Tests:** Automatically rolls back after each test
 - **@Rollback:** Control rollback behavior in tests
 - **@Commit:** Force commit in test methods
 - **Test Isolation:** Ensure tests don't interfere with each other
-

8. Common Pitfalls and Troubleshooting

Common Issues

1. **Self-Invocation:** Internal method calls bypass transaction proxy
2. **Private Methods:** `@Transactional` ignored on non-public methods
3. **Exception Handling:** Catching exceptions prevents rollback
4. **Connection Leaks:** Long-running transactions can exhaust connection pool
5. **Deadlocks:** Poor transaction design can cause database deadlocks

Debugging Tips

- **Enable Transaction Logging:** Set appropriate log levels
- **Monitor Connection Pool:** Watch for connection exhaustion
- **Profile Performance:** Identify slow transactions
- **Test Rollback Scenarios:** Verify exception handling works correctly

Monitoring and Metrics

- **Transaction Duration:** Monitor transaction execution time
- **Rollback Frequency:** Track rollback rates
- **Connection Usage:** Monitor database connection utilization
- **Lock Contention:** Identify database locking issues